

Microservices

Patterns . Tools . Technologies

Krishna Mohan Koyya
August 2018

Microservices

Patterns

Solution Architecture

Problem Statement

- Server-side enterprise application
- Variety of different clients
 - desktop browsers, mobile browsers, native mobile applications
- Northbound API for 3rd parties
 - REST API
- Enterprise Integration
 - Web services and/or message broker.
- HTML/JSON/XML responses
- *What's the application's deployment architecture?*

Forces

- Understandability and Maintainability
- Quick response to frequent business requirements: Agile
- 24/7 Availability and High Scalability
- Team dynamics: New faces, rotation
- Technology Adaptation

Monolithic Architecture

Problem Statement

- Server-side enterprise application
- Variety of different clients
 - desktop browsers, mobile browsers, native mobile applications
- Northbound API for 3rd parties
 - REST API
- Enterprise Integration
 - Web services and/or message broker.
- HTML/JSON/XML responses
- What's the application's deployment architecture?

Forces

- Understandability and Maintainability
- Quick response to frequent business requirements: Agile
- 24/7 Availability and High Scalability
- Team dynamics: New faces, rotation
- Technology Adaptation

Solution

- Tiered and Layered Architecture: UI+Service+Data, MVC, DAO
- One single distribution: WAR

Negative Impact

- Loaded IDE, Delayed deployments, Build Nightmares, Outdated Technologies

Microservice Architecture

Problem Statement

- Server-side enterprise application
- Variety of different clients
 - desktop browsers, mobile browsers, native mobile applications
- Northbound API for 3rd parties
 - REST API
- Enterprise Integration
 - Web services and/or message broker.
- HTML/JSON/XML responses
- What's the application's deployment architecture?

Forces

- Understandability and Maintainability
- Quick response to frequent business requirements: Agile
- 24/7 Availability and High Scalability
- Team dynamics: New faces, rotation
- Technology Adaptation

Solution

- Application as a set of loosely coupled collaborating services
- Polyglot databases and technologies
- Self-contained deployments
- Focussed Teams with CI/CD governance

Negative Impact

- Complexity associated with any distributed systems

Decomposition

Problem Statement

- Microservice Architecture
- Small, agile, autonomous and cross-functional teams
- Continuous Delivery and Deployment
- Independent Deployment
- Simplified Testing
- Polyglot Databases and Technology Stacks
- *How to decompose the application into services?*

Forces

- Architectural stability
- Loosely coupled and cohesive services
- Single Responsibility Principle
- Common Closure Principle

Decomposition by Business Capability

Problem Statement

- Microservice Architecture
- Small, agile, autonomous and cross-functional teams
- Continuous Delivery and Deployment
- Independent Deployment
- Simplified Testing
- Polyglot Databases and Technology Stacks
- *How to decompose the application into services?*

Forces

- Architectural stability
- Loosely coupled and cohesive services
- Single Responsibility Principle
- Common Closure Principle

Solution

- Discover services on the lines of business capability
- Core and Expandable Architecture
- Organization Structure
- High-Level Domain Model

Challenges

- Very difficult to discover services: There is no single way

Decomposition by Subdomains

Problem Statement

- Microservice Architecture
- Small, agile, autonomous and cross-functional teams
- Continuous Delivery and Deployment
- Independent Deployment
- Simplified Testing
- Polyglot Databases and Technology Stacks
- *How to decompose the application into services?*

Forces

- Architectural stability
- Loosely coupled and cohesive services
- Single Responsibility Principle
- Common Closure Principle

Solution

- Discover services on the lines of Domain Driven Design
- Core Subdomains: In-house development
- Supporting Subdomains: Outsourcing
- Generic Subdomains: Off the shelf solutions

Challenges

- Very difficult to discover services: Requires strong understanding of the business.

Deployment

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

One Host - Many Service Instances

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

Solution

- Multiple instances of service per host
 - Multiple JVMs on a host
 - One Tomcat per JVM
 - One or more services per Tomcat
- Leads to better utilisation of the host resources

Negative Impact

- Conflicting resource requirements
- Conflicting dependency versions
- Difficult demarcation
- Difficulty in monitoring individual services

One Host - One Service Instance

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

Solution

- Single instance of service per host
 - Only one JVM on a host
 - One Tomcat per JVM
 - One service per Tomcat
- No conflict in resource requirements
- No dependency conflict
- Full control on demarcation
- Difficulty in monitoring individual services

Negative Impact

- Less efficient utilisation of the host resources

One VM - One Service Instance

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

Solution

- Single instance of service per VM like EC2 instance
 - Only one JVM on a VM
 - One Tomcat per JVM
 - One service per Tomcat
- Services are isolated ... better demarcation and no conflicts
- IaaS case: unmatched scaling

Negative Impact

- Building VM image is slow and time consuming

One Container - One Service Instance

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

Solution

- Single instance of service per container
 - Only one JVM on a container
 - One Tomcat per JVM
 - One service per Tomcat
- Services are isolated ... better demarcation and no conflicts
- Very good for scaling
- Easy deployment and relatively less load on the host

Negative Impact

- Relatively newer technology

Serverless Deployment

Problem Statement

- Microservice architecture
- Smaller services
- Multiple instances for throughput
- Multiple instances for availability
- *How to package and deploy the services?*

Forces

- Heterogenous set of languages, frameworks & versions
- Multiple instances of each service
- Independent deployment
- Resource consumption constraints
- Monitoring
- Quick and frequent deployment

Solution

- Public clouds with opaque and elastic infrastructure
- Deploy only the functionality
 - AWS Lambda, GC Funtions, Azure Functions
- No heavy lifting of image creation and deployment

Negative Impact

- Limited possibilities in terms of supported technologies
- Limited integration capabilities beyond REST API and MQTT
- Elastic, means no pre-provisioning
- Bad for long-duration services
- Vendor-locking

External API

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Multiple kinds of clients
- Client needs several services
- *How to expose the services to client?*

Forces

- Granularity of the API
- Different data needs from the same service
- Mobile Vs WAN Vs LAN performance
- Number and location of services varies
- Service partitioning/aggregation over a period of time
- Change in protocols

API Gateway

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Multiple kinds of clients
- Client needs several services
- *How to expose the services to client?*

Forces

- Granularity of the API
- Different data needs from the same service
- Mobile Vs WAN Vs LAN performance
- Number and location of services varies
- Service partitioning/aggregation over a period of time
- Change in protocols

Solution

- API Gateway as the single point of entry
- All call reach API Gateway
- The API Gateway proxies the requests to the actual services
- Event-driven mechanism (Node/Vertx/Reactor)

Negative Impact

- Impact of changes in service on the single API Gate way
- One-size-fits-all may not really true
- Latency

Backend to Frontend

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Multiple kinds of clients
- Client needs several services
- *How to expose the services to client?*

Forces

- Granularity of the API
- Different data needs from the same service
- Mobile Vs WAN Vs LAN performance
- Number and location of services varies
- Service partitioning/aggregation over a period of time
- Change in protocols

Solution

- Different API Gateways for different contexts (like client types)
- Specific calls reach specific API Gateway
- Specific API Gateway proxies the requests to the actual services
- Event-driven mechanism (Vertx/Node/Reactor)

Negative Impact

- Impact of changes in service on each of the API Gate ways
- Latency

Security

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Multiple instances of the services
- *How to secure the service?*

Forces

- Stateless calls
- No server-side sessions
- Ever-changing granularity of the services
- Client capabilities
- Social Networks

Access Tokens

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Multiple instances of the services
- *How to secure the service?*

Forces

- Stateless calls
- No server-side sessions
- Ever-changing granularity of the services
- Client capabilities
- Social Networks

Solution

- Access Tokens
- Client sends access token in every call
- Authentication at the gateway
- JSON Web Tokens or JWT
- Open Authentication

Cross Cutting Concerns

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Configuration
- Logging
- Monitoring
- Health Checks
- *How to deal with the cross cutting concerns across the application?*

Forces

- Quick development
- Optimal investment on something common across services

Chasis

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Configuration
- Logging
- Monitoring
- Health Checks
- *How to deal with the cross cutting concerns across the application?*

Forces

- Quick development
- Optimal investment on something common across services

Solution

- Build a chassis
- Spring Boot Actuator
 - /actuator
- External Log Storage, Log Aggregation, Correlation IDs
 - AWS Cloudwatch
 - LogStash

Externalisation

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Configuration
- Logging
- Monitoring
- Health Checks
- *How to deal with the cross cutting concerns across the application?*

Forces

- Quick development
- Optimal investment on something common across services

Solution

- Build a chassis
- Spring Boot Actuator
 - /actuator
- External Log Storage, Log Aggregation, Correlation IDs
 - AWS Cloudwatch
 - LogStash
- External Configuration
 - Docker Compose supplies the configuration as OS variables

Collaboration

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Collaboration
- *How two microservices of the application communicate with each other?*

Forces

- Frequent Communication
- Heavy payloads
- Binary payloads
- Synchronous and Asynchronous
- Coupling

Remote Procedure Invocation

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Collaboration
- *How two microservices of the application communicate with each other?*

Forces

- Frequent Communication
- Heavy payloads
- Binary payloads
- Synchronous and Asynchronous
- Coupling

Solution

- Use Request/Reply RPI calls
- REST: HTTP/1
- gRPC: HTTP/2 with duplex communication
- Apache Thrift with binary payloads

Negative Impact

- No support for notifications, pub/sub
- Reduced availability: If one goes down, the other goes down
- Difficulty associated with discovering other service

Messaging

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Collaboration
- *How two microservices of the application communicate with each other?*

Forces

- Frequent Communication
- Heavy payloads
- Binary payloads
- Synchronous and Asynchronous
- Coupling

Solution

- Use messaging for inter-service communication
- Data-pipelining with Apache Kafka: publish/subscribe, fault-tolerant, stream processing
- Rabbit MQ: Publish/subscribe

Negative Impact

- Additional Complexity with yet another technology
- Clients need to discover message broker
- General complexity associated with asynchronous communication

Data Management

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- What is the database architecture?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Shared Database

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- What is the database architecture?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Solution

- Use single database server for the whole application
- Stick to ACID transactions

Negative Impact

- Development-time coupling
- Run-time coupling
- Chosen database may not satisfy the storage requirements of some of the services
- Scaling issues

Database per Service

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- What is the database architecture?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Solution

- Each service have it's own database
- Enforce demarcation, may be using different user-ids for different services
 - Tables per service
 - Schema per service
 - Database Server per service

Negative Impact

- Cross-service transaction are still complex
- Joining data across services is still challenging

Command-Query Segregation

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- What is the database architecture?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Solution

- Let the services of the application divide based on the responsibility
 - Services that insert/delete/update records (commanding service)
 - Services that query the records (querying service)
- Let the commanding service fire data-change events
- Let the querying service subscribe to the events and update their views
- De-normalized databases

Negative Impact

- Increased complexity
- Replication lag
- Eventual consistency

Event Sourcing

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- What is the database architecture?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Solution

- Let the commands result in only state-change events
- Let the queries build the state of the entity by replaying the events
- Inherently atomic
- Use it with command-query segregation
- Provision of snapshots ... to reduce amount replay
- Suitable for event driven systems
- Temporal audit

Negative Impact

- Relatively newer technology
- Increased complexity

API Composition

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Common data domain
- Database per service
- What is the query join strategy?

Forces

- Services must be loosely coupled
- Business transactions spanning across services
- Queries across services
- Data joining across several services
- Replication of data for scaling
- Different storage requirements

Solution

- Let an API composer join the query results in-memory

Negative Impact

- Very in-efficient

Service Discovery

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- How a client service discover a provider service?

Forces

- Each instance of a service runs at a specific location (host & Port)
- Number of services/instance and locations change dynamically
- Containers are assigned with dynamic IP addresses

Service Registry

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- How a client service discover a provider service?

Forces

- Each instance of a service runs at a specific location (host & Port)
- Number of services/instance and locations change dynamically
- Containers are assigned with dynamic IP addresses

Solution

- Implement a naming registry of the services
- Register a service location against a name at the service startup
- Deregister the service location at the service shutdown
- Registry polls the services time to time
- Several readymade offerings
 - Netflix Eureka with Spring Boot
 - Consul
 - etcd
 - Apache ZooKeeper

Challenges

- Registry must be highly available
- Fixed location of the registry service must be known to the clients

Self Registration

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How a client registers with the registry?

Forces

- Instances must be registers at startup
- Instance must be deregistered at shutdown
- Crashed instances must be removed from the registry
- Incapable instances must be identified and removed from the registry

Solution

- Instances registers and deregisters themselves
- Instances periodically renew the registration (heart-beats)
- Instances can provide fine granular states (STARTING, STARTED, PAUSED ...)
- Build as part of chassis

Challenges

- Coupling between service and registry
- Incapable service may not aware the reality
- Chassis for various platforms

Registrar

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How a client registers with the registry?

Forces

- Instances must be registers at startup
- Instance must be deregistered at shutdown
- Crashed instances must be removed from the registry
- Incapable instances must be identified and removed from the registry

Solution

- A 3rd party registrar takes note of start/shutdown of services and updates the register
- A service composer takes this responsibility
 - Eg. Kubernetes
 - Eg. Netflix Prana for non-JVM applications
- Individual services are free from working with the registers
- Build as part of chassis

Challenges

- Registrar can have only superficial knowledge of the state of the service
- Adds one more crucial component to the infrastructure

Client-side Discovery

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How a client discover the service?

Forces

- Each instance of a service runs at a specific location (host & Port)
- Number of services/instance and locations change dynamically
- Containers are assigned with dynamic IP addresses
- Registry is having the info of the services

Solution

- Let the client queries the registry to find the service
- Lesser number of variables

Challenges

- Apart from the services, the clients are also coupled with the registry
- Each language should have the chassis ... for all the micro services

Server-side Discovery

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How a client discover the service?

Forces

- Each instance of a service runs at a specific location (host & Port)
- Number of services/instance and locations change dynamically
- Containers are assigned with dynamic IP addresses
- Registry is having the info of the services

Solution

- Let the load balancer query the registry to find the service, on behalf of the client
- No coupling between the client and registry

Challenges

- More moving parts
- More protocols to be supported by the API Gateway
- Each language should have the chassis ... for all the micro services

Reliability

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How to mitigate if a service is unavailable?
- How to prevent cascading affect?

Forces

- A depending service may be unavailable
- A depending service may be taking too much time
- A client retrying for the failed service consumes resources

Circuit Breaker

Problem Statement

- Microservice architecture
- Smaller and multiple services
- Service collaboration
- Service Registry is implemented
- How to mitigate if a service is unavailable?
- How to prevent cascading affect?

Forces

- A depending service may be unavailable
- A depending service may be taking too much time
- A client retrying for the failed service consumes resources

Solution

- Wrap the calls to the service with a monitor
- The monitor works like electric circuit breaker
- When the monitor is CLOSED, the calls are passed to the service
- When the monitor is OPEN, the calls returned to the client with error
- When the monitor is HALF-OPEN, the calls attempted on the service
 - If service responds, monitor CLOSES, otherwise OPENS
- Timeout/threshold to move from CLOSED to OPEN
- Timeout to move from OPEN to HALF-OPEN
- Netflix Hystrix

Challenges

- Choosing the timeouts/thresholds
- Yet another cog in the wheel

Microservices

Tools and Technologies

Spring Boot

Eureka . Hystrix . Ribbon . Circuit Breaker

Microservices

Tools and Technologies

Docker Containers

Microservices

Tools and Technologies

Docker Compose

Microservices

Tools and Technologies

CI/CD

Git. Jenkins . Pipelines . Dockers

Thanks

Glarimy Technology Services, Bengaluru, India

www.glarimy.com | <https://github.com/glarimy>

Krishna Mohan Koyya

krishna@glarimy.com